

Seminar Advances in Deep Learning for Time Series (ADLTS)

Lecture 3: Deep Learning for Time Series

Dr. Dario Zanca

Machine Learning and Data Analytics (MaD) Lab
Friedrich-Alexander-Universität Erlangen-Nürnberg
17.10.2024

Team: FAU & PUCV



Dr. Dario Zanca (FAU)
dario.zanca@fau.de



Naga Venkata Sai Jitin Jami, M. Sc. (FAU)
jitin.jami@fau.de



Dr. Christoffer Loeffler (PUCV)
christoffer.loffler@pucv.cl



&



PONTIFICIA
UNIVERSIDAD
CATÓLICA DE
VALPARAÍSO


- I. Introduction
 - II. The Tool Tracking dataset
 - III. DL for Time Series
 - IV. Time-aware models
 - V. XAI for Time Series - part 1
 - VI. Active Learning for Time Series - part 1
 - VII. Semi-supervised Learning
 - VIII. Domain-shifts, Ethics, and Bias
 - IX. XAI for Time Series - part 2
 - X. Active Learning for Time Series - part 2
-

- I. Introduction
 - II. The Tool Tracking dataset
 - III. DL for Time Series**
 - IV. Time-aware models
 - V. XAI for Time Series - part 1
 - VI. Active Learning for Time Series - part 1
 - VII. Semi-supervised Learning
 - VIII. Domain-shifts, Ethics, and Bias
 - IX. XAI for Time Series - part 2
 - X. Active Learning for Time Series - part 2
-

References

Deep Learning

by Ian Goodfellow, Yoshua Bengio, and Aaron Courville (2016)



Deep Learning: Foundations and Concepts

by C. Bishop, H. Bishop (2024)



Lecture outline

1. Introduction to Deep Learning
 2. Convolutional Neural Networks (CNNs)
 3. Recurrent models (RNNs and LSTMs)
 4. Transformers
-



ADLTS \ DL for TS \

Introduction to Deep Learning



Why Deep Learning?

Previous method needed **handcrafted features**:

- MFCCs (speech processing) (1)
- I-Vector (speech processing) (2)
- Sift (scene alignment, videos) (3) → Needs expert knowledge about domain

(1) Mermelstein, P. (1976). Distance measures for speech recognition, psychological and instrumental. *Pattern recognition and artificial intelligence*, 116, 374-388.

(2) V. Gupta, P. Kenny, P. Ouellet and T. Stafylakis, "I-vector-based speaker adaptation of deep neural networks for French broadcast audio transcription," *2014 IEEE International Conference on Acoustics, Speech and Signal Processing (ICASSP)*, 2014, pp. 6334-6338, doi: 10.1109/ICASSP.2014.6854823.

(3) Lowe, D. G. (2004). Distinctive image features from scale-invariant keypoints. *International journal of computer vision*, 60(2), 91-110.

Why Deep Learning?

What if we can not define generally applicable features?

- High dimensional data
- Hard to come up with generally applicable features

→ With deep learning, we can find features in a data driven way

→ Can help capture complex non-linear relationships

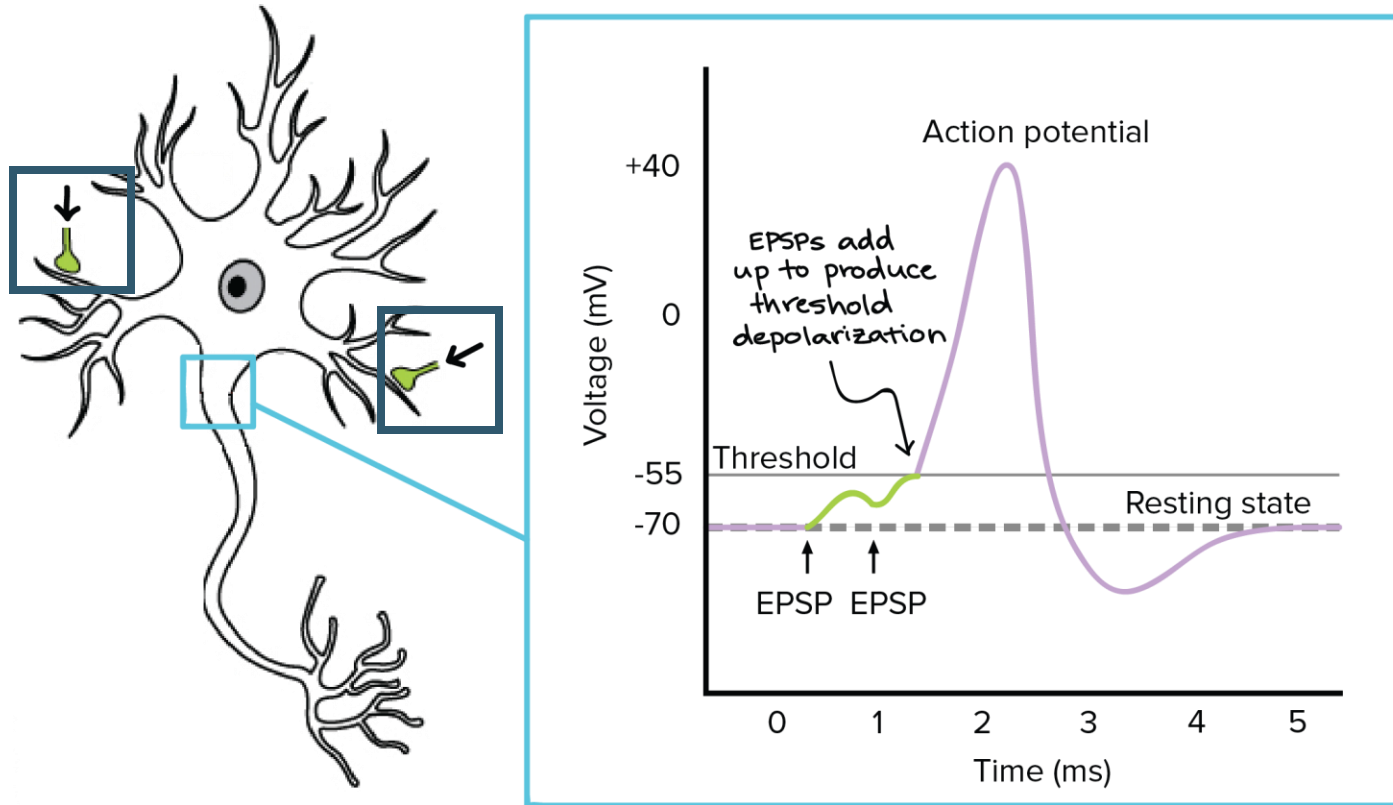
The Human Brain

The human brain is our reference for an intelligent agent.

- It contains different areas specialized for some tasks (e.g., the **visual cortex**)
- It consists of neurons as the fundamental unit of “**computation**”



The Brain's Neuron



EPSP = Excitatory postsynaptic potential

- Excitatory **stimuli** reach the neuron
- Threshold is reached
- Neuron fires and triggers **action potential**

The Perceptron – Computational Model of a Neuron

Let's build the computational model step by step:

1. Show the **input** and **output** of our neuron (which depends on the data and task)

Input x_1



Input x_2



Input x_3




Output y




The Perceptron – Computational Model of a Neuron

2. **Weights** can “select” or “deselect” **input** channels (not all are relevant for subsequent computations)




The Perceptron – Computational Model of a Neuron

3. We add up all the excitatory signals and the **resting potential (or bias)** to determine the current potential.



The Perceptron – Computational Model of a Neuron

4. A **threshold function (or activation function) σ** is applied to determine whether an action potential has to be sent in the **output**



Activation Functions

The goal of an **activation function** is to activate (or not) neuron.


It brings **non-linearity** to a neural network:

→ The output is **not** a linear function of the input

Different activation functions can be considered, taking into account their:

- Output range
- Mean
- Gradient

Sigmoid




$$\sigma(x) = \frac{1}{1 + e^{-x}}$$

The Perceptron – Computational Model of a Neuron

5. We can write the perceptron mathematical model to map **inputs** x_1, x_2, x_3 to the **output** y_1 using channel **weights** w_1, w_2, w_3, b :

$$y = \sigma(w_1 \cdot x_1 + w_2 \cdot x_2 + w_3 \cdot x_3 + b) = \sigma\left(\sum_i w_i \cdot x_i + b\right)$$




Single Layer Perceptron

We can combine multiple perceptrons to create a **layer**.

One perceptron has the following equation:

$$y_1 = \sigma(w_{11} \cdot x_1 + w_{12} \cdot x_2 + w_{13} \cdot x_3 + b)$$




Single Layer Perceptron

We can combine multiple perceptrons to create a **layer**.

We can combine **three perceptrons** into one layer:

$$\begin{pmatrix} y_1 \\ y_2 \\ y_3 \end{pmatrix} = \underbrace{\begin{pmatrix} w_{11} & w_{12} & w_{13} \\ w_{21} & w_{22} & w_{23} \\ w_{31} & w_{32} & w_{33} \end{pmatrix}}_W \begin{pmatrix} x_1 \\ x_2 \\ x_3 \end{pmatrix} + \underbrace{\begin{pmatrix} b_1 \\ b_2 \\ b_3 \end{pmatrix}}_b$$



Or in a more simplified form: $y = \sigma(W \cdot x + b)$

Multi Layer Perceptron (MLP)

We can chain **multiple layers**


$$y = \sigma(W^1 \cdot x + b^1)$$

$$z = \sigma(W^2 \cdot y + b^2)$$

So: $z = \sigma(W^2 \cdot \sigma(W^1 \cdot x + b^1) + b^2)$

→ Each layer has its own set of parameters (weights W^i and bias b^i)

→ The activation function ensures that all layers do not collapse into one: it introduces **non-linearity**




Multilayer perceptron (MLP)

We call “**hidden layer**” any layer in between the input and the output layers.

For example, this neural network (image on the right) is a Multi-Layer-Perceptron (MLP) with a **single hidden layer** (highlighted by the green box).


→ The underlying computation is described by:
$$\mathbf{y}^{i+1} = \sigma(\mathbf{W}^i \cdot \mathbf{y}^i + \mathbf{b}^i)$$

Also referred to feedforward networks (no feedback connection)



How are Model Parameters Learned?

- 1) **Forward propagation:** This phase refers to the computation of the output using input and parameters.
- 2) **Loss calculation:** The output and expected output are then compare using a loss function.
- 3) **Backward propagation:** During this phase, the model computes the gradients of the loss with respect to each parameter (θ).



Backpropagation Algorithm: Overview

1. Calculate the forward pass and **store** results for \hat{y} , a_j^k , and z_j^k .
 2. Calculate the backward pass and **store** results for $\frac{\partial L}{\partial w_{ij}}$, proceeding from the last layer:
 - a) **Evaluate** the error terms for the last layer δ_k
 - b) **Backpropagate** the error term for the computation of δ_j
 - c) **Iterate** to all previous layers
 3. Combine the individual gradients (average)
 4. Update the weights according to the **learning rate** λ
-




ADLTS \ DL for TS \ Convolutional Neural Networks (CNNs)



Motivation: The Human Visual Cortex



CNNs were inspired by a specialized brain area: the **visual cortex**.

- 1) The visual cortex processes visual information collected by the retinae in a **hierarchical** manner.
- 2) The **receptive field** of cortical neurons increases the later they are in this hierarchy.
 - a) Small receptive fields are stimulated by high spatial frequencies - or fine details
 - b) Large receptive fields are stimulated by low spatial frequencies - or coarse details



Modeling Receptive Fields

- In standard **multilayer perceptrons** (MLP), layers are **fully-connected**, i.e. all units from one layer are connected to all units from the previous layer.
 - Cortical neurons have small **receptive fields**: units from one layer have **sparse** and **localized** connection with units from the previous layer.
- This is modeled by discrete convolutional operations.



Convolutional Operation


The **continuous convolution** is a operation on two functions f and g noted $f * g$ that produces a third function.

$$(f * g)(t) = \int_{-\infty}^{\infty} f(\tau)g(t - \tau)d\tau$$

It is the integral of the product of the two functions after one is reflected about the y-axis and shifted.

In deep learning, a **discrete** version of this operation is used.

Image from Pihlajamäki, Tapani. Multi-resolution Short-time Fourier Transform Implementation of Directional Audio Coding.



Discrete Convolution

The (discrete) **convolution** is a mathematical operation on two functions. On our case, the **input** and a smaller **filter** produce a third function (the feature map).

$$(input * filter)[n] = \sum_{m=-\infty}^{+\infty} input[m] \cdot filter[n - m]$$


→ The sum is evaluated for all values of the shift.

→ The term convolution is used both to indicate the result function and process of computing it.

1-D Convolutional Neural Networks

CNNs can also be used to learn **temporal dependencies** on time series data.

- The convolution is applied along a single dimension, i.e., the **temporal** dimension.
- The resulting model is generally referred to as 1-D CNN.



1D Convolutional Operation

Like the 2D operation, the output computation only depends on a **subset** of the time series:


$$y_1 = w_1 x_1 + w_2 x_2 + w_3 x_3$$



1D Convolutional Operation

Like the 2D operation, the output computation only depends on a **subset** of the time series:


$$y_2 = w_1 x_2 + w_2 x_3 + w_3 x_4$$



1D Convolutional Operation

Like the 2D operation, the output computation only depends on a **subset** of the time series:

$$y_3 = w_1 x_3 + w_2 x_4 + w_3 x_5$$



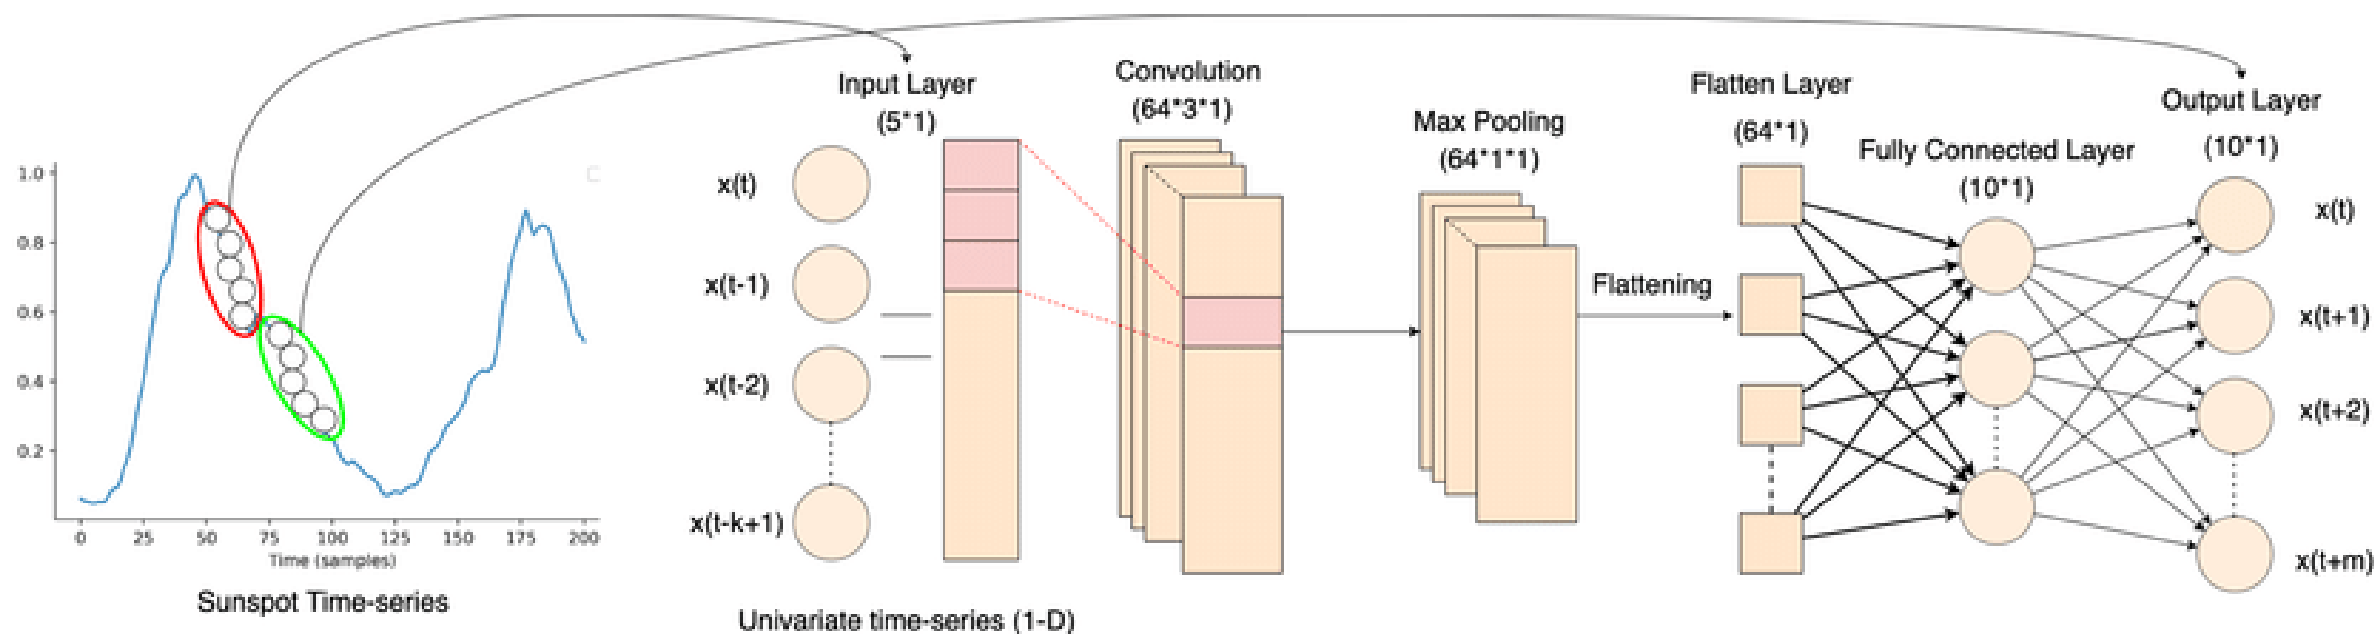
CNN Properties and Time Series

CNN properties make them very appropriate for **computer vision**. Those can also be useful for **time series**:

- Ability to extract **local patterns**: 1D CNNs can extract local patterns that occur in time series data (e.g. short term trends in stock prices, weather events)
 - **Translation Equivariance**: 1D CNNs can identify those patterns regardless of their position, which is useful for tasks such as anomaly detection
 - **Dimensionality reduction** of pooling layers: This can be useful when datasets are high dimensional.
 - **Hierarchical** feature learning: Time series are composed of patterns at different time **scales**. (e.g. seasonal weather patterns vs sudden weather event)
-

CNN for Time Series Forecasting

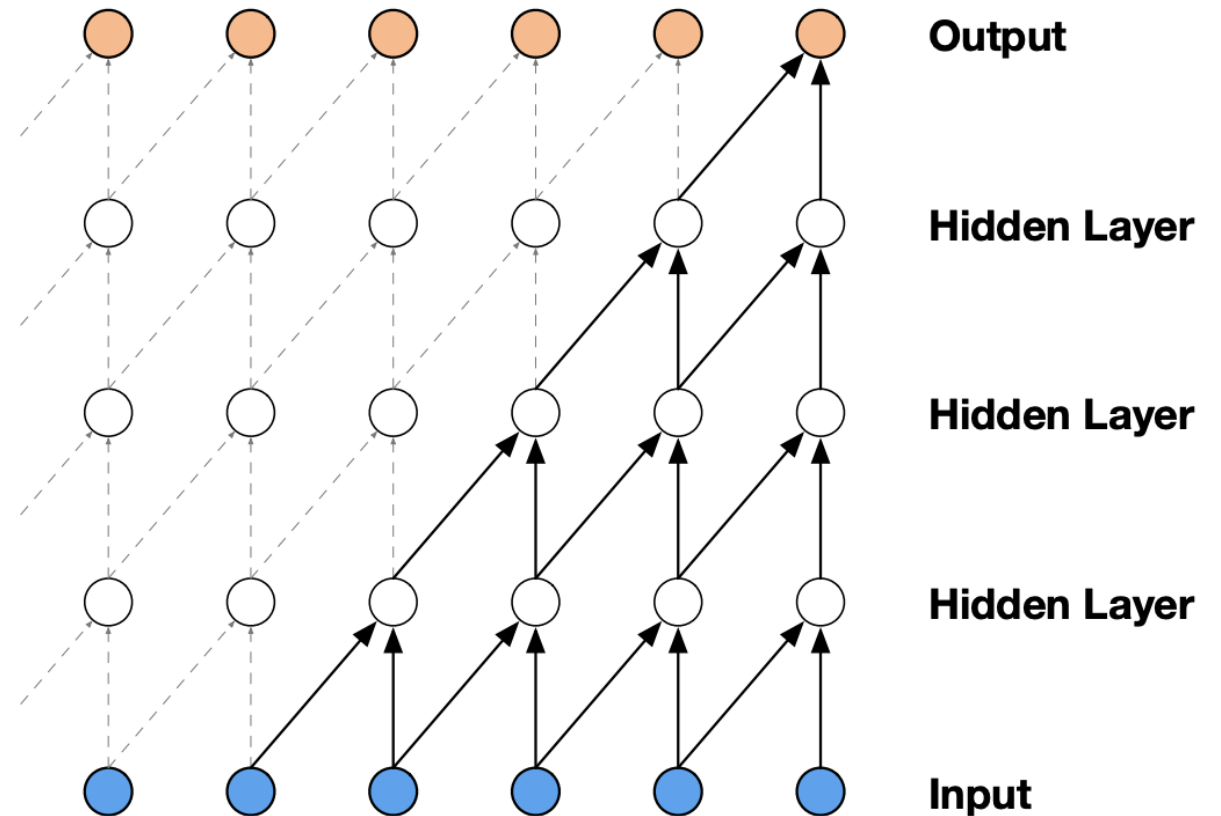
Example of CNN for univariate time series **forecasting** [1].



Temporal Convolutional Networks (TCNs)

The **causal convolution** is best suited to model causality in the data.

- For images it can be implemented by “masked convolutions”, i.e., a tensor mask is applied before the actual convolution takes place.
- For 1D data, e.g., audio processing, it can be more easily implemented by shifting the output of a normal convolution by a few timesteps.





ADLTS \ DL for TS \

Recurrent models (RNNs and LSTMs)




Limitations of NN for time series data


Feed-forward and Convolutional neural networks present some **disadvantages**, when applied to sequential data:

1. Cannot work online (sequence has to be fed all at once)
2. Consider only the current input
 - Cannot memorize previous time steps
3. Cannot handle directly sequences of different lengths


Recurrent Neural Network (RNN)




Recurrent Neural Network (RNN)



Recurrent Neural Network (RNN)



Recurrent Neural Network (RNN)



A **recurrent neural network (RNN)** is a neural network that contains feed-back connections.

- Activations can flow in a loop
- It allows for temporal processing

An RNN is composed by:

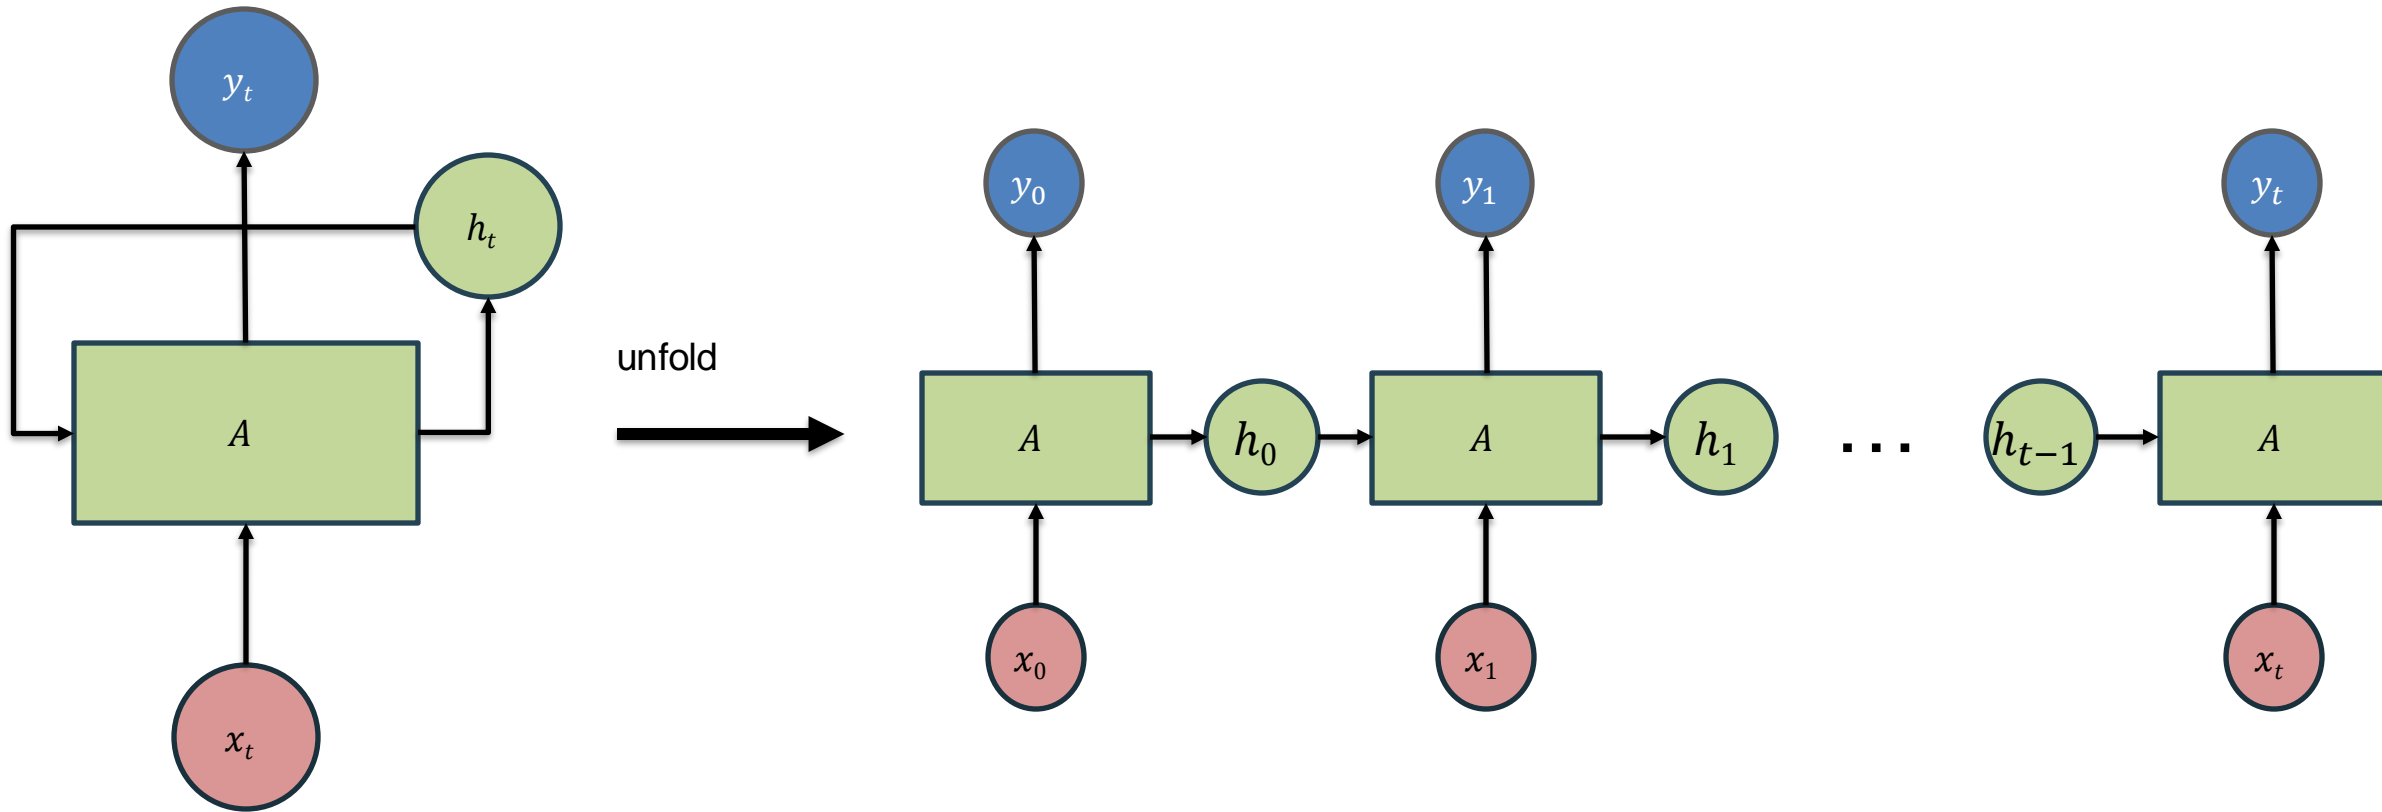
x_t : input at time t

A : neural network


h_t : hidden state

y_t : output at time t


RNN Unfolding



Mathematical formulation



Mathematical formulation




The behaviour of the RNN can be described as a **dynamical system** by the pair of non-linear matrix equations:

$$h_t = \tanh(W_{hh}h_{t-1} + W_{xh}x_t)$$


$$y_t = \sigma(W_{hy}h_t)$$

The order of the dynamical system corresponds to the dimensionality of the state h_t .

RNNs architecture




Example: one to many



A typical example of a one to many problem is that of **image captioning**.

Input:



Output:

A

cat


playing

with

a

ball

Example: many to one



A typical example of a many to one problem is that of **sentiment analysis**.

Input:

Horrible

service

the

room

was


dirty

Output:



("negative")

Example: many to many



A typical example of a many to many problem is that of name entity recognition.


Input:

Harry Potter and Hermione invented a new spell

Output:

1 1 0 1 0 0 0 0

Example: many to many



Another example of a many to many problem is that of **machine translation**.

Input:

Horrible

service

the

room

was

dirty

Output:

Un

servizio

orribile

la

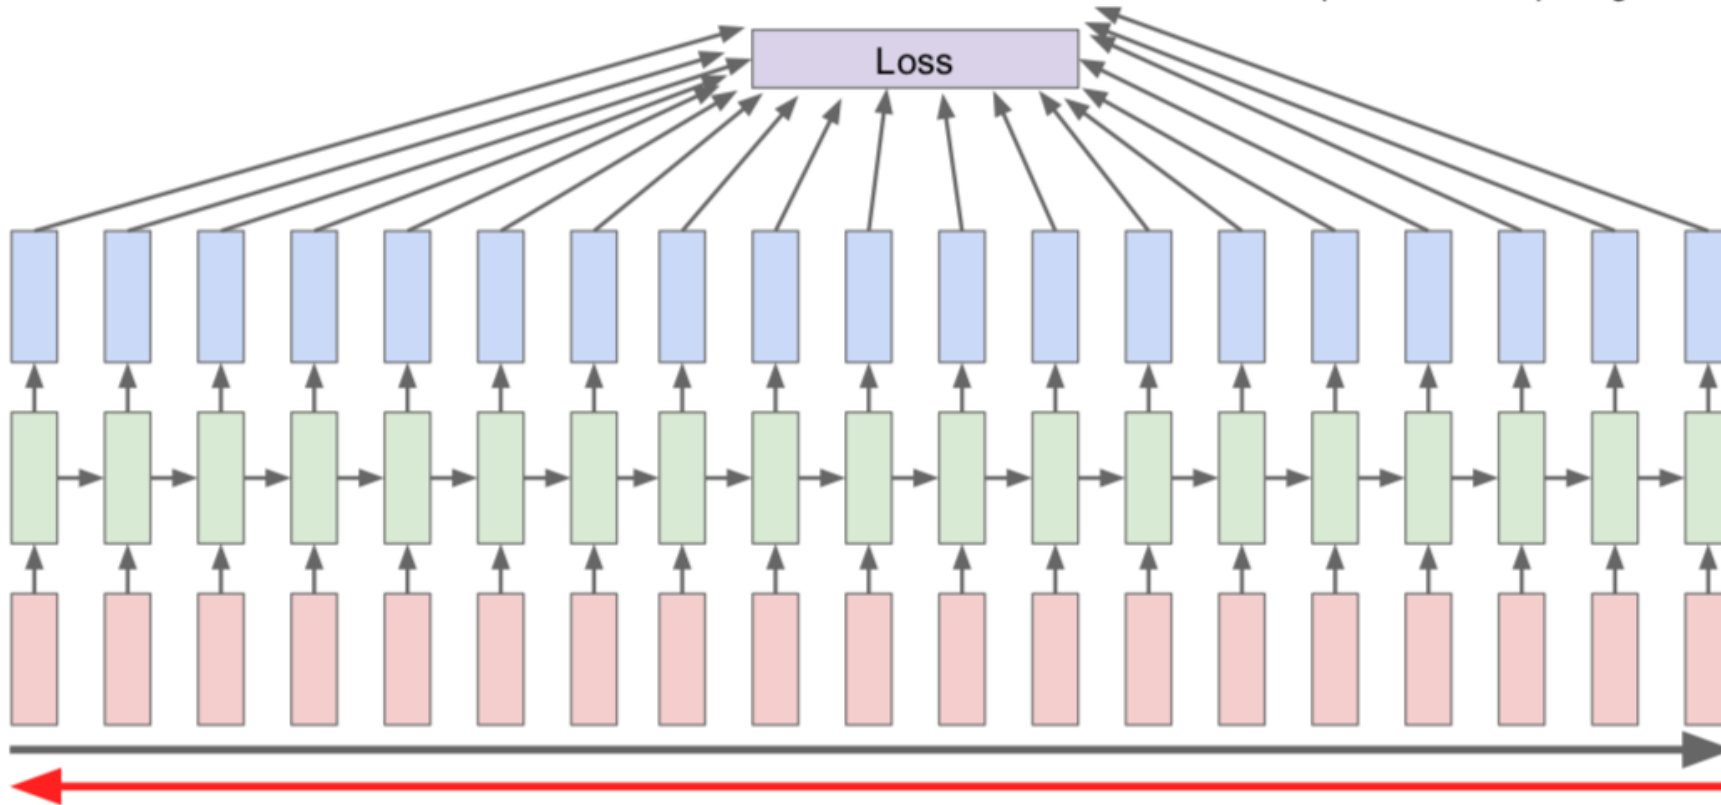
camera

era

sporca


BPTT: Limitations

BPTT can be computationally very expensive as a lot of partial derivatives have to be computed, depending on the complexity of the network.




Truncated Backpropagation Through Time (Trunc-BPTT)


With the Truncated Backpropagation Through Time (**Trunc-BPTT**), instead of passing the whole sequence, we perform the forward and backpass on a subset.




Long-short term memory networks (LSTMs)



Long-short term memory networks (LSTMs)




Long-short term memory networks (LSTMs)



Observations:

- Compared to RNNs, additionally to the **hidden state** h_t , we have another state, c_t , said the **cell state**.


Long-short term memory networks (LSTMs)



Observations:

- Compared to RNNs, additionally to the **hidden state** h_t , we have another state, c_t , said the **cell state**.
- The information flow is regulated by **three gates**: the forget gate, the input gate and the output gate.


Long-short term memory networks (LSTMs)



The cell state has:

- Only minor interactions
- Simple information flow
- Other gates regulates whether it is preserved/not-preserved or updated.


Long-short term memory networks (LSTMs)



The **forget gate** decides how much information to retain from the previous cell state.

$$f_t = \sigma(W_f \cdot [h_{t-1}, x_t] + b_f)$$

Long-short term memory networks (LSTMs)




The **input gate** decides the information to be added to the cell state, based on the current input.

$$i_t = \sigma(W_i \cdot [h_{t-1}, x_t] + b_i)$$

$$g_t = \tanh(W_c \cdot [h_{t-1}, x_t] + b_c)$$


Long-short term memory networks (LSTMs)



The values can be combined to update the cell state

$$c_t = f_t \odot c_{t-1} + i_t \odot g_t$$

Long-short term memory networks (LSTMs)




The **output gate** generates the output of the current LSTM cell, based on the current input, the previous output, and the updated cell state.

$$o_t = \sigma(W_o \cdot [h_{t-1}, x_t] + b_o)$$

$$h_t = o_t * \tanh(c_t)$$

Long-short term memory networks (LSTMs)



To summarize, the LSTM cell is described by the following equations:

$$f_t = \sigma(W_f \cdot [h_{t-1}, x_t] + b_f)$$

$$i_t = \sigma(W_i \cdot [h_{t-1}, x_t] + b_i)$$

$$g_t = \tanh(W_c \cdot [h_{t-1}, x_t] + b_c)$$

$$c_t = f_t \odot c_{t-1} + i_t \odot g_t$$

$$o_t = \sigma(W_o \cdot [h_{t-1}, x_t] + b_o)$$

$$h_t = o_t * \tanh(c_t)$$



ADLTS \ DL for TS \ Transformers



Attention Is All You Need [1]

The **Transformer** [1] architecture was published in 2017. Models built on this architecture have become state-of-the-art in many domains, starting with **Natural Language Processing**.

- Completely built on the **self-attention** mechanism
- Does **not** use sequence any recurrent architecture:
 - More efficient: input sequences can be processed in **parallel**
 - Has no inherent understanding of sequence order
- Can be applied to various tasks, including in time series machine learning



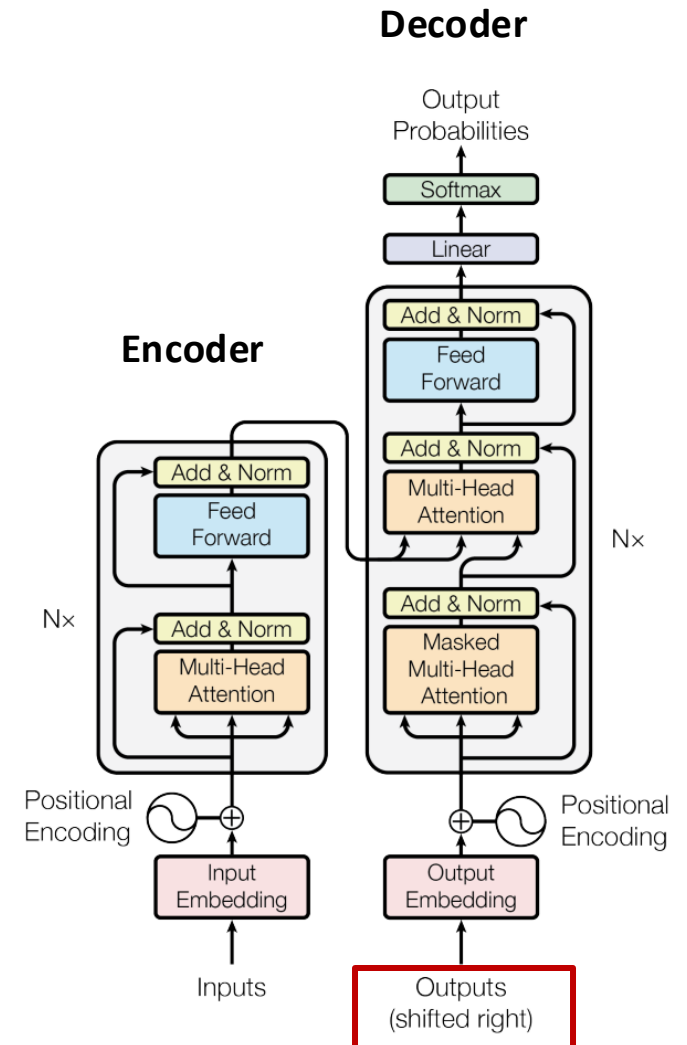
The Transformer Architecture

The fundamental components of the Transformer architecture are:

- **Positional encoding**
- **Multi-head self attention** based scaled dot product attention
- An **encoder-decoder** architecture

The transformer was first proposed as a **machine translation** model.

“Attention is all you need”, Vaswani, et al.



Transformer: Positional Encoding

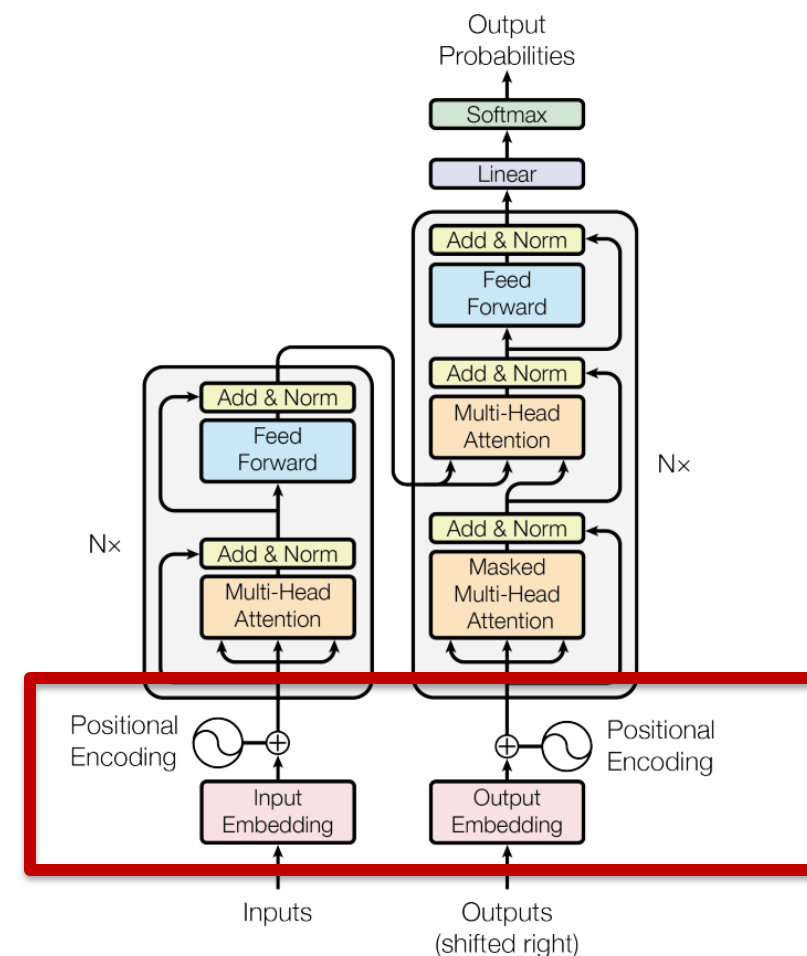
Positional encoding is added to the input and output sequence embeddings in the encoder and decoder. It is necessary to give information on **word order**.

Positional encoding is defined by using sine and cosine functions with different frequencies:


$$PE_{(pos, 2i)} = \sin\left(\frac{pos}{10000^{\frac{2i}{d_{model}}}}\right)$$

$$PE_{(pos, 2i+1)} = \cos\left(\frac{pos}{10000^{\frac{2i}{d_{model}}}}\right)$$

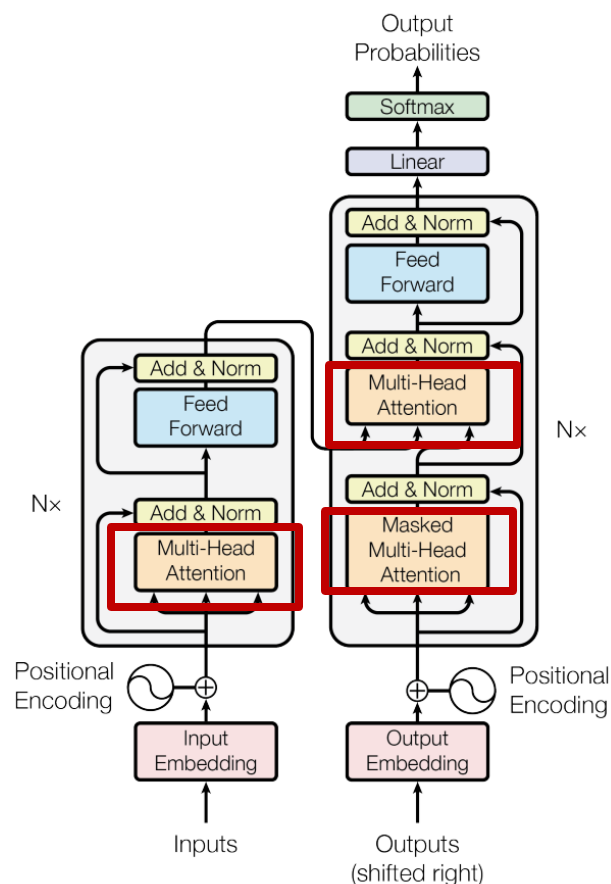
where pos is the position and i is the embedding dimension.



Transformer: Positional Encoding



Transformer: Multi-head Attention Mechanisms



Three multi-head attention mechanisms are used in the Transformer architecture:


- **Encoder self-attention:** captures relationships between input sequence tokens
- **Decoder self-attention:** captures relationships between output sequence tokens
- **Encoder-decoder cross-attention:** captures relationships between input and output sequences

Transformer: Multi-Head Self-Attention

Multi-head attention is used to compute attention several times in **parallel**, using **independent** weight matrices $W^{V,i}$, $W^{K,i}$ and $W^{Q,i}$. These operations are called **attention heads**, and are then **concatenated**.

Basic scaled dot product attention is not sufficient to encode the **complexity** of language, as it might only focus on one aspect of relationships between tokens.

For instance, the multi-head attention mechanism can focus both on **long range and short range relationships** through different attention heads.



Multi-head attention

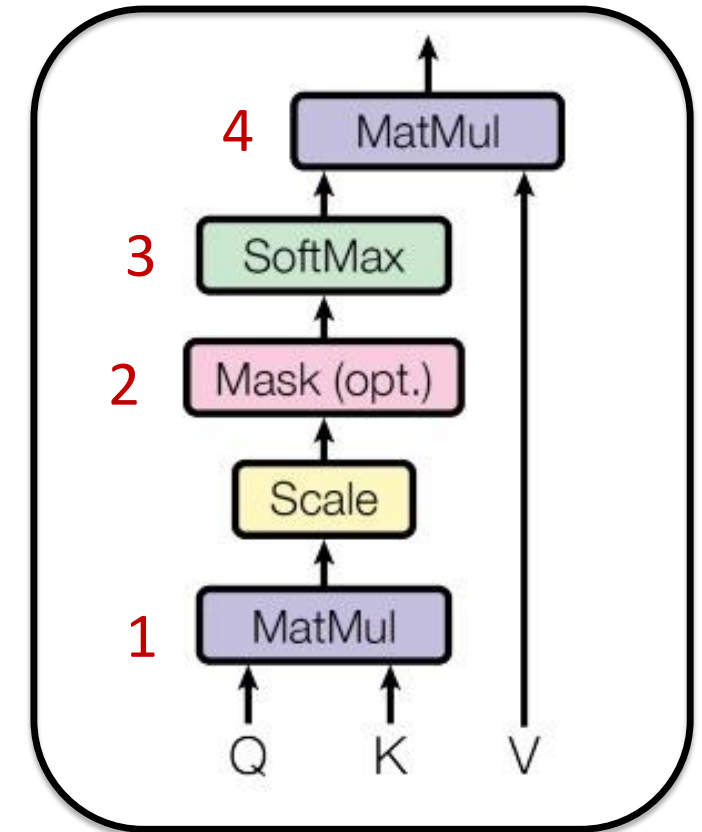
Transformer: Scaled Dot Product Attention

Then, the scaled dot-product attention is computed by:

$$\text{Attention}(Q, K, V) = \text{softmax}\left(\frac{QK^T}{\sqrt{d_k}}\right)V$$

where d_k is the dimension of Q and K .

1. How each token is related to the query (i.e., similarity score)
2. Converts into probabilities (sum to 1)
3. **Causal masking** is used to prevent information leakage
4. Weighted sum of values that highlights relevant vectors

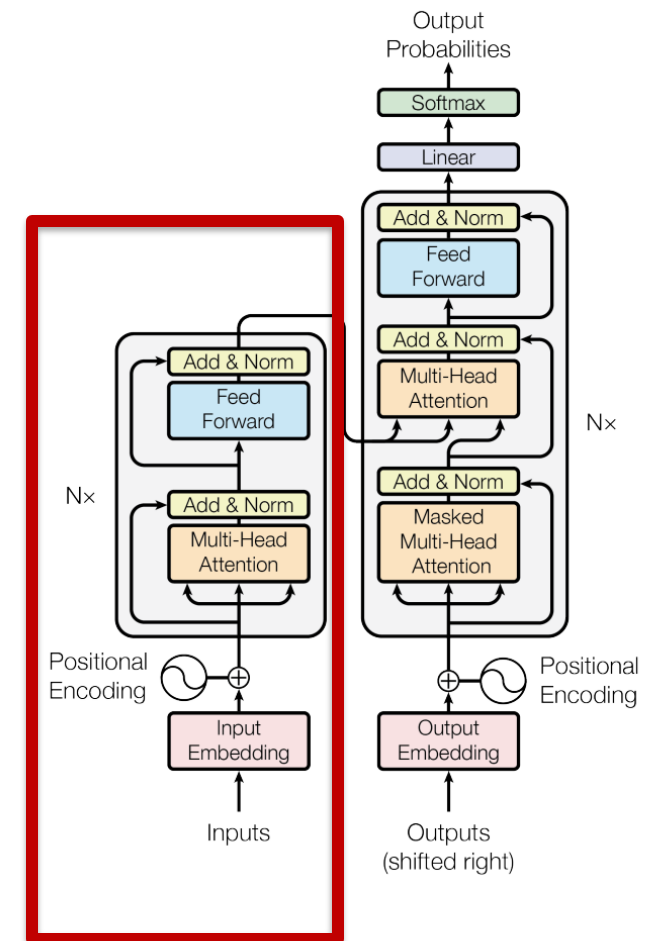


Scaled Dot-Product Attention

Transformer: Encoder

The encoder generates a **contextualized representation** of the **input sequence**. This representation has the same length as the input sequence, and is used to condition the decoder.

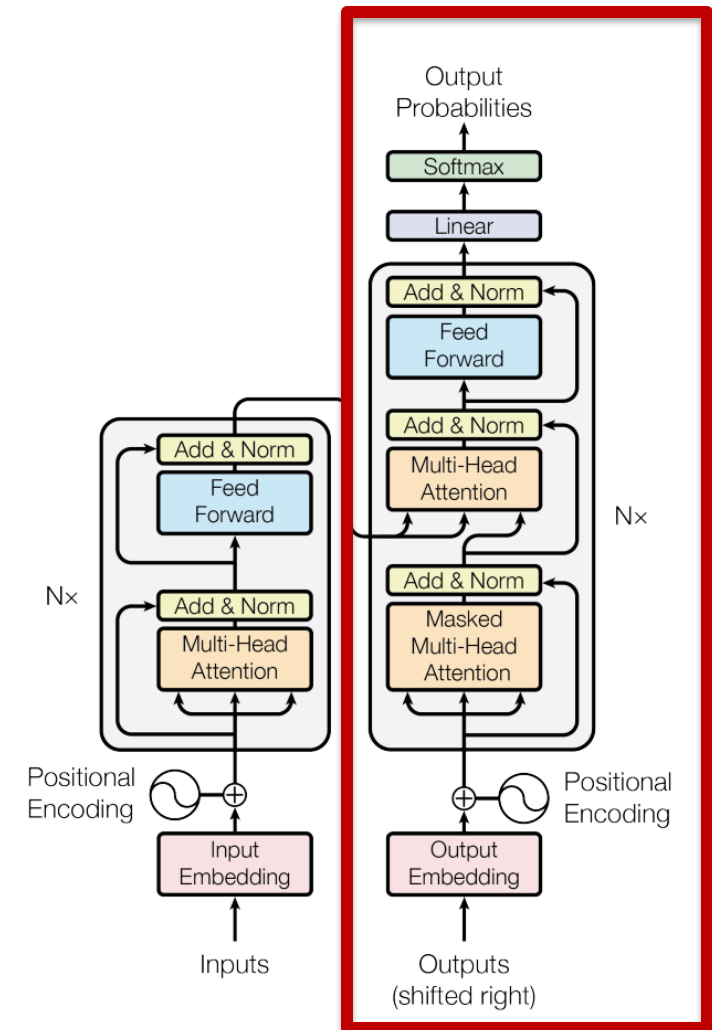
- It is made of N_x identical layers, each composed of
 - Multi-head **self-attention**
 - Feedforward neural network
 - Residual connection
 - Layer normalization



Transformer: Decoder

The **decoder** generates the output of the model (e.g., translated text). It generates tokens one at a time, using network information from the **encoded representation**.

- It is also made of N_x identical layers:
 - **Masked** multi-head self-attention
 - Multi-head **cross-attention**
 - Feedforward neural network
 - Residual connection
 - Layer normalization



Transformer Model: Pros and Cons

The main advantages of Transformer models are:

- **Parallelization:** Transformers can process all tokens in a sequence simultaneously.
- **Long-Range Dependencies:** Transformers can capturing dependencies between distant tokens, as long as they are all present in the same sequence

However:

- **Quadratic time complexity** of attention mechanism
-



ADLTS \ DL for TS \ Recap



Lecture outline

1. Introduction to Deep Learning
 2. Convolutional Neural Networks (CNNs)
 3. Recurrent models (RNNs and LSTMs)
 4. Transformers
-



